



ECE 552: Introduction to Computer Architecture

15. Pipelining, Hazards & Exceptions #3

Lecture notes based in part on slides created by Joshua San Miguel, Mikko Lipasti, Mark Hill, David Wood, Guri Sohi, John Shen, and Jim Smith



Refresher 1/3

❑ Pipeline function for ADD (R-type):

- Fetch: read instruction from memory
- Decode: read source operands from reg
- Execute: calculate sum
- Memory: pass results to next stage
- Writeback: write sum into register file

```
add x1, x2, x3 # write to x1  
add x4, x1, x5 # read from x1
```

* RAW data dependency



Refresher 2/3

- ❑ Data Dependency: *one instruction uses the result of a previous one*
 - ❑ Doesn't necessarily cause a problem (e.g., WAR, WAW)
- ❑ Data Hazard: *one instruction has a data dependency that will cause a problem if we don't "deal with it"*



Refresher 3/3

Avoid

➤ Make sure there are no hazards in the code

Detect and Stall

➤ If hazards exist, stall the processor until they go away.

Detect and Forward

➤ If hazards exist, fix up the pipeline to get the correct value (if possible)



Exercise

lw x1, 0(x2)

lw x3, 4(x1)

add x5, x4, x3

- (a) Assume **no RF bypassing and no forwarding**
- (b) Assume **RF bypassing but no forwarding**
- (c) Assume **RF bypassing and forwarding**

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								

Exercise

lw x1, 0(x2)
 lw x3, 4(x1)
 add x5, x4, x3

(a) No RF bypassing and no forwarding

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
lw x3			F	D*	D*	D*	D	X	M	W			
add				F*	F*	F*	F	D*	D*	D*	D	X	M

Exercise

lw x1, 0(x2)

lw x3, 4(x1)

add x5, x4, x3

(b) RF bypassing but no forwarding

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
lw x3			F	D*	D*	D	X	M	W				
add				F*	F*	F	D*	D*	D	X	M	W	

Exercise

```

lw x1, 0(x2) # Result available after the MEM stage
lw x3, 4(x1) # The result of I1 is needed during the EX stage
add x5, x4, x3

```

(c) RF bypassing and forwarding

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
lw x3			F	D*	D	X	M	W					
add				F*	F	D*	D	X	M	W			

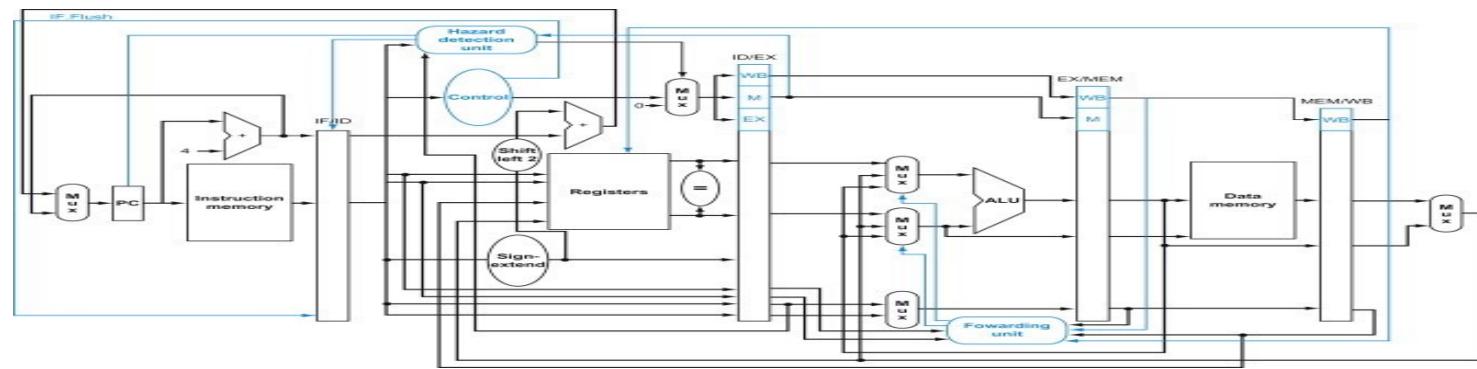
MEM-to-EX forwarding

MEM-to-EX forwarding

Exercise

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
nop				X	M	W							
lw x3		F	D*	D	X	M	W						
nop						X	M	W					
add			F*	F	D*	D	X	M	W				

cycle 3:



add

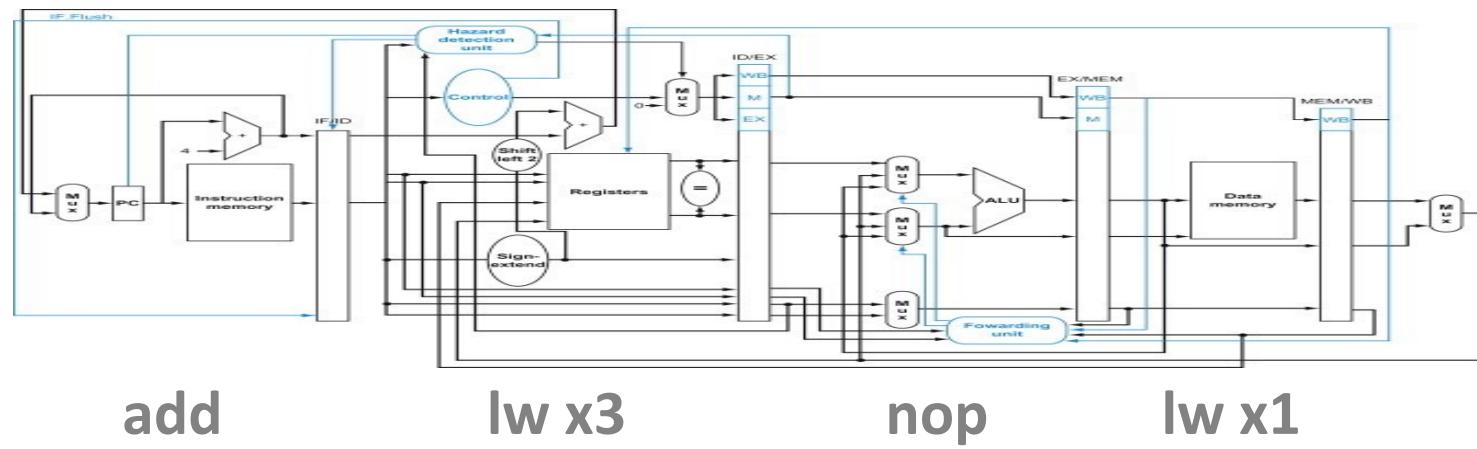
lw x3

lw x1

Exercise

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
nop				X	M	W							
lw x3		F	D*	D	X	M	W						
nop						X	M	W					
add			F*	F	D*	D	X	M	W				

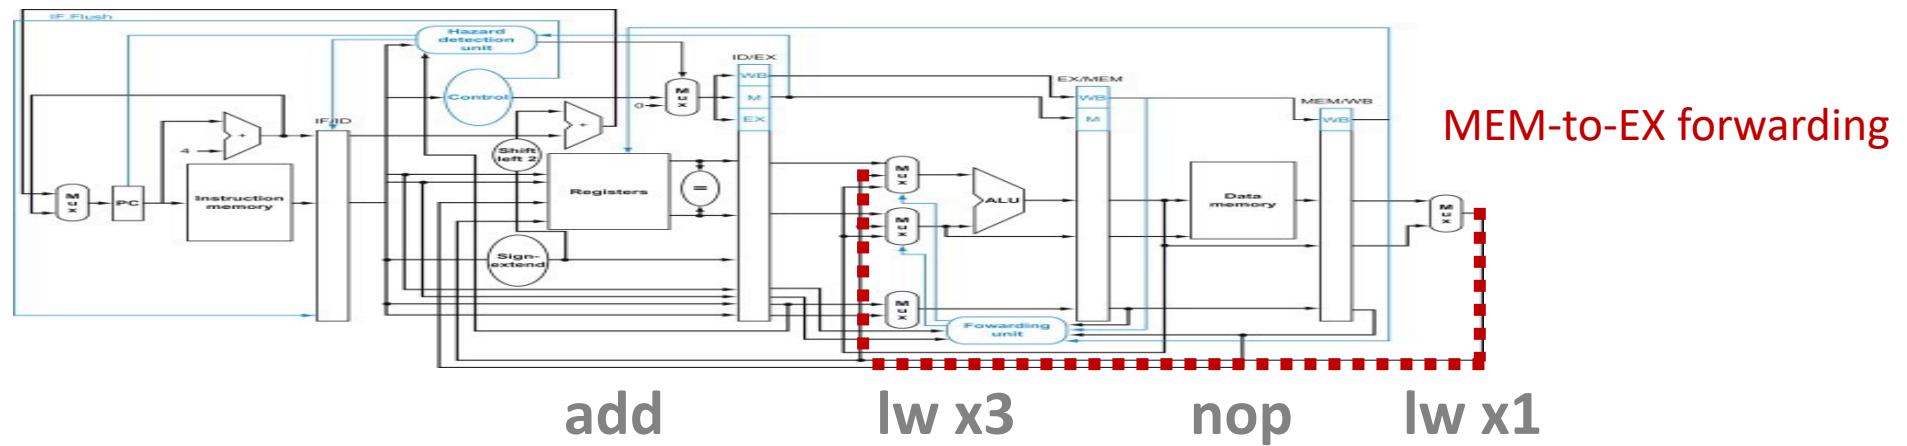
cycle 4:



Exercise

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
nop					X	M	W						
lw x3			F	D*	D	X	M	W					
nop							X	M	W				
add				F*	F	D*	D	X	M	W			

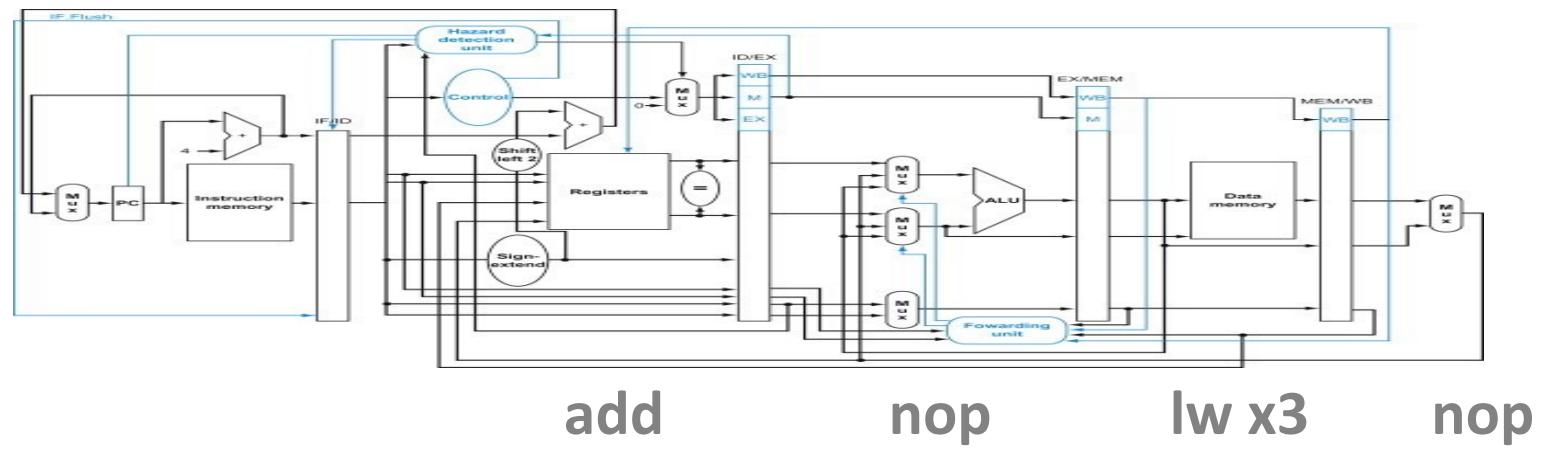
cycle 5:



Exercise

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
nop					X	M	W						
lw x3			F	D*	D	X	M	W					
nop							X	M	W				
add				F*	F	D*	D	X	M	W			

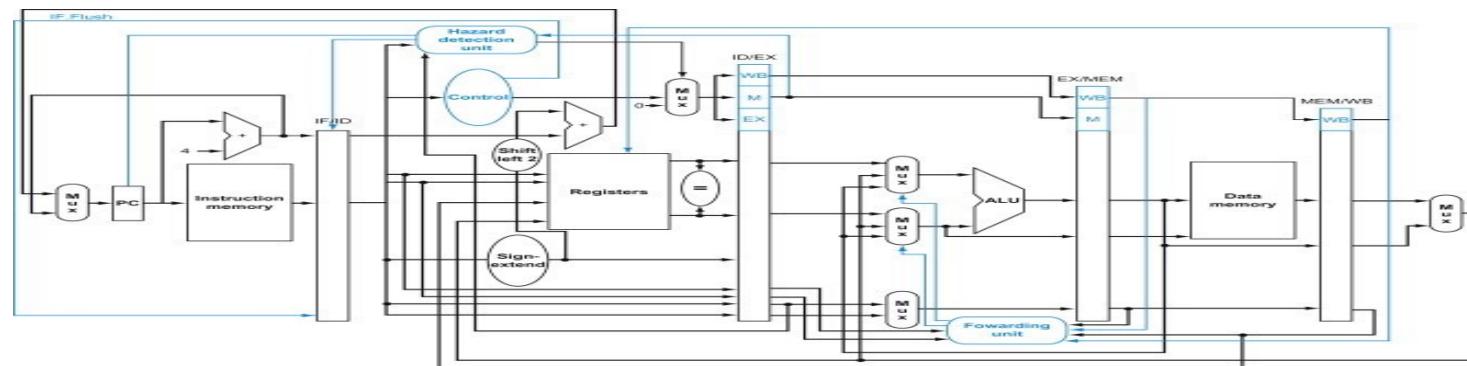
cycle 6:



Exercise

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw x1	F	D	X	M	W								
nop					X	M	W						
lw x3			F	D*	D	X	M	W					
nop							X	M	W				
add				F*	F	D*	D	X	M	W			

cycle 7:



add

nop

lw x3



One more exercise

(a) Assume **RF bypassing but no forwarding**

(b) Assume **RF bypassing and forwarding**

* Annotate all RF bypassing and forwarding as appropriate

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
add x1, x2, x3	F	D	X	M	W											
nor x3, x4, x5																
add x6, x3, x7																
lw x3, 0(x6)																
sw x6, 4(x3)																

One more exercise

(a) Assume **RF bypassing** but **no forwarding**

* Annotate all RF bypassing and forwarding as appropriate

	inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
WAR	add x1, x2, x3	F	D	X	M	W											
RAW	nor x3, x4, x5			F	D	X	M	W									
RAW	add x6, x3, x7				F	D*	D*	D	X	M	W						
WAR	lw x3, 0(x6)					F*	F*	F	D*	D*	D	X	M	W			
RAR	sw x6, 4(x3)							F*	F*	F	D*	D*	D	X	M	W	
RAW																	



One more exercise

(b) Assume **RF bypassing** and **forwarding**

* Annotate all RF bypassing and forwarding as appropriate



Control dependencies

- An instruction Y is control-dependent on an instruction X if the outcome of X determines whether or not Y will be executed
 - Due to conditional branches:
 - Branch must execute to determine which instruction to fetch next
 - Instructions following a conditional branch are control-dependent on the branch instruction

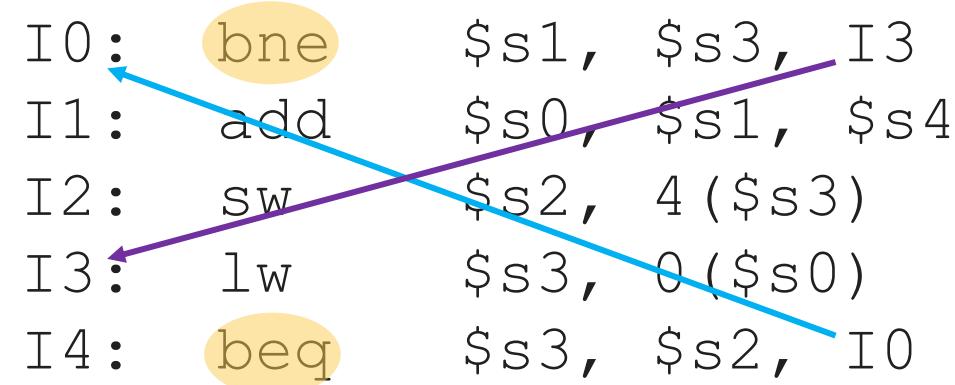


Control dependencies

I0:	bne	\$s1, \$s3, I3
I1:	add	\$s0, \$s1, \$s4
I2:	sw	\$s2, 4(\$s3)
I3:	lw	\$s3, 0(\$s0)
I4:	beq	\$s3, \$s2, I0

- ❑ List all control dependences

Control dependencies



❑ List all control dependences:

- I0→I1, I0→I2
- I4→I0, I4→I1, I4→I2, I4→I3, I4→I4

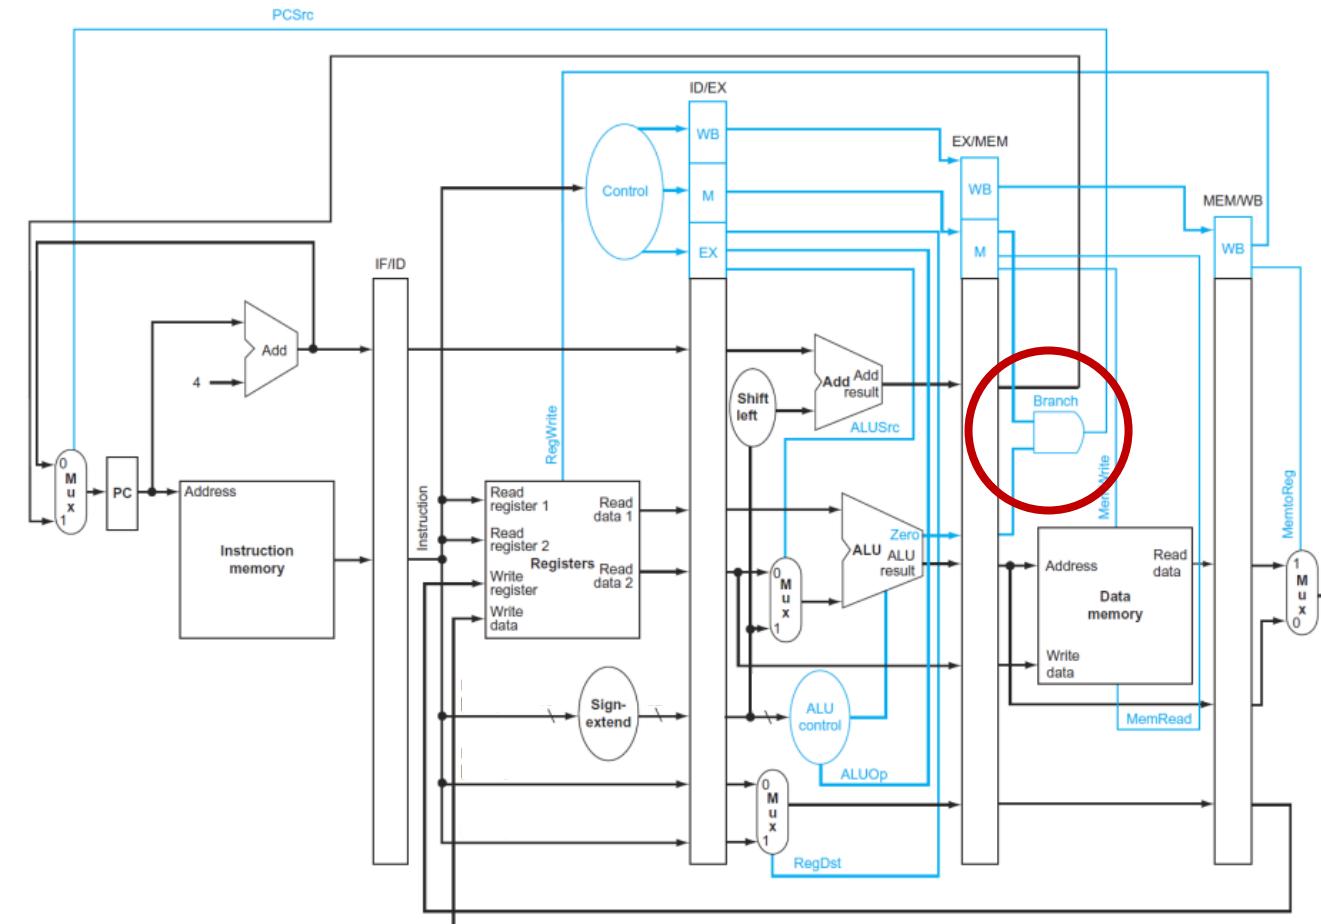


Control dependencies

- Technically, control dependences need not be preserved
 - Can execute instructions that should not have been executed, as long as program correctness is not affected:
 - Data dependences are preserved
 - No new exceptions incurred

Pipeline control hazards

- ❑ When an incorrect instruction is fetched because the outcome of a previous branch has not yet been resolved
 - e.g., Assume branches resolved in M stage



Pipeline control hazards

- ❑ When an incorrect instruction is fetched because the outcome of a previous branch has not yet been resolved
 - e.g., Assume branches resolved in M stage

beq x0, x1, TARGET
 lw x3, 0(x0)
 add x4, x5, x6
 or x2, x1, x7

TARGET:
 sub x5, x0, x5

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
beq	F	D	X	M	W								
lw			F	D	X	M	W						
add				F	D	X	M	W					
or					F	D	X	M	W				
sub						F	D	X	M	W			



Pipeline control hazards

- When an incorrect instruction is fetched because the outcome of a previous branch has not yet been resolved
 - e.g., Assume branches resolved in M stage

Predict-not-taken:

- Fetch next instructions *speculatively*
- flush if branch is actually taken

beq x0, x1, TARGET
lw x3, 0(x0)
add x4, x5, x6
or x2, x1, x7

TARGET:
sub x5, x0, x5

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
beq	F	D	X	M	W								
lw			F	D	X	M	W						
add				F	D	X	M	W					
or					F	D	X	M	W				
sub						F	D	X	M	W			

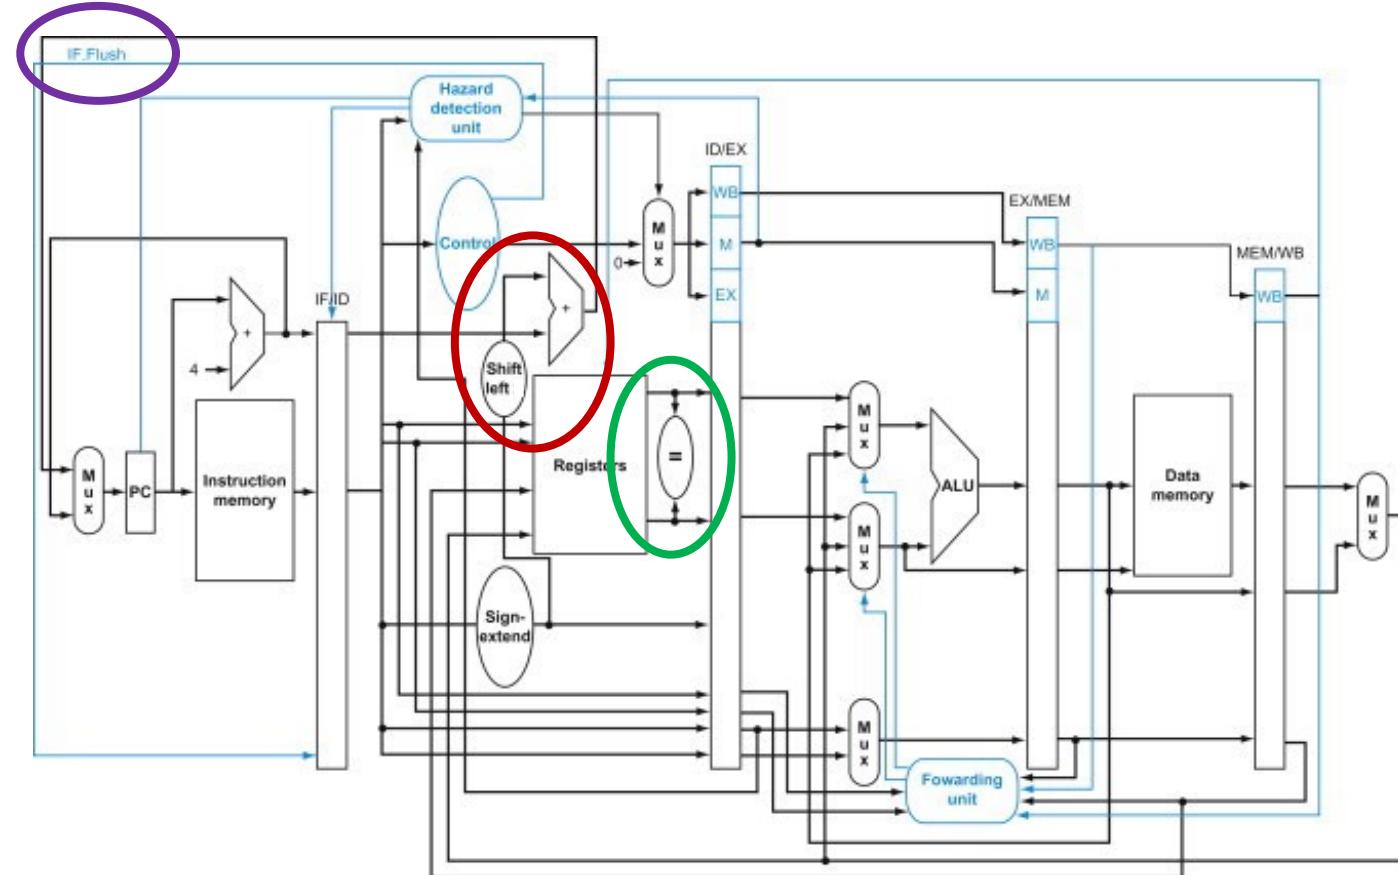
Resolving control hazards

□ Predict-not-taken + resolve in ID:

- Compute target and compare in ID stage to reduce branch penalty
- Flush: clear IF/ID.Instruction

2. Compare register operands as soon as possible
3. Clear IF/ID.Instruction

1. Compute the branch target PC in the decode stage



Pipeline control hazards

- When an incorrect instruction is fetched because the outcome of a previous branch has not yet been resolved
 - e.g., Assume branches resolved in D stage

beq x0, x1, TARGET
lw x3, 0(x0)
add x4, x5, x6
or x2, x1, x7

TARGET:
sub x5, x0, x5

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
beq	F	D											
lw			F										

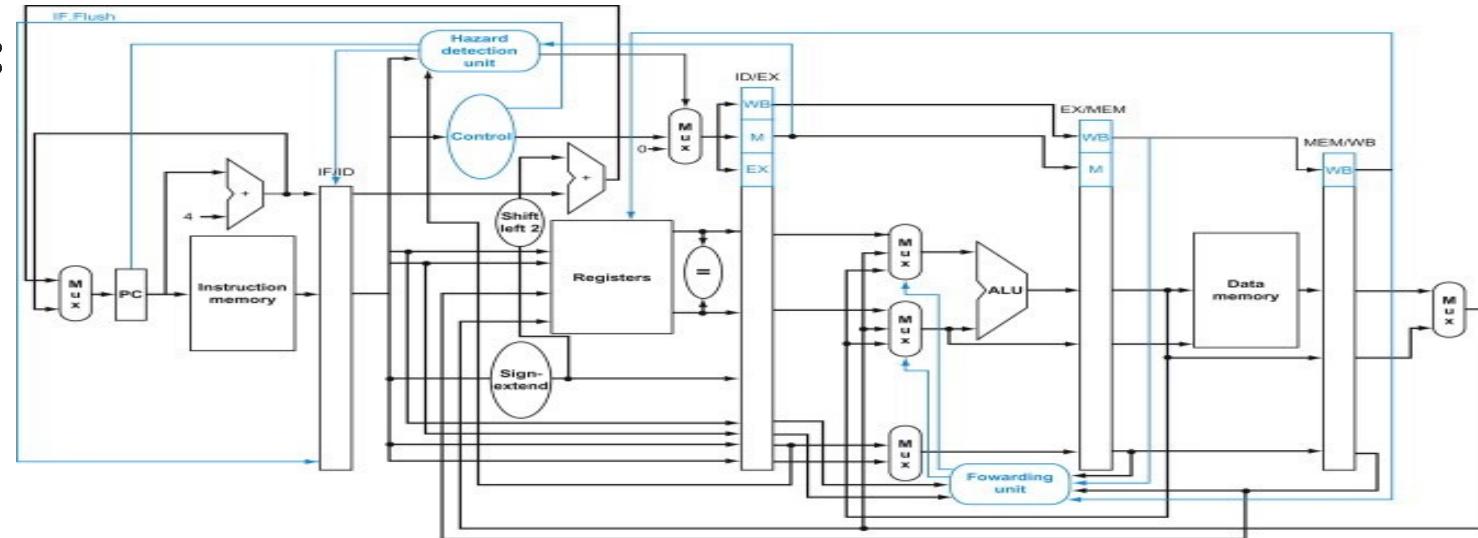
* Flush 1 cycle, not 3 (if needed)

Pipeline control hazards

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
beq	F	D	X	M	W								
lw			F	=									
sub				F	D	X	M	W					

beq x0, x1, TARGET
 lw x3, 0(x0)
 add x4, x5, x6
 or x2, x1, x7
 TARGET:
 sub x5, x0, x5

cycle 8:



sub

lw (NQF)

lw (NQF)

lw (NQF)

lw (NQF)

Pipeline control + data hazards

- ❑ Predict-not-taken + resolve in ID:
 - Compute target and compare in ID stage
 - Flush: clear IF/ID.Instruction
 - **But RAW hazards in ID stage?**
 - Stall + EX-to-ID forwarding

add x1, x5, x6
beq x0, x1, TARGET

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
add	F	D	X	M	W								
beq			F	D*	D	X	M	W					



Branch delay slots

□ Branch delay slot:

- Instruction after branch *always* executed (avoids flush logic)
 - Compiler is responsible for filling slot with *legal* instruction
- Assume branches resolved in ID stage

lw x7, 0(x8)
beq x0, x1, TARGET
add x5, x4, x6

TARGET:
sub x5, x0, x5



inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw	F	D	X	M	W								
beq		F	D	X	M	W							
add			F	D	X	M	W						
sub				F	D	X	M	W					



Branch delay slots

□ Branch delay slot:

- Instruction after branch *always* executed (avoids flush logic)
 - Compiler is responsible for filling slot with *legal* instruction
- Assume branches resolved in ID stage

lw x7, 0(x8)

beq x0, x1, TARGET

nop

add x5, x4, x6

TARGET:

sub x5, x0, x5

Option 1:

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
lw	F	D	X	M	W								
beq		F	D	X	M	W							
nop													
sub					F	D	X	M	W				



Branch delay slots

□ Branch delay slot:

- Instruction after branch *always* executed (avoids flush logic)
 - Compiler is responsible for filling slot with *legal* instruction
- Assume branches resolved in ID stage

~~lw x7, 0(x8)~~

beq x0, x1, TARGET

TARGET:

sub x5, x0, x5

Option 2:

lw x7, 0(x8)

add x5, x4, x6

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
beq	F	D	X	M	W								
lw			F	D	X	M	W						
sub				F	D	X	M	W					



Branch prediction

- Dynamic branch direction prediction:
 - *Predict* branch outcome (T or N) in IF stage
 - Assume branches resolved in ID stage

beq \$t0, \$t1, TARGET
add \$t5, \$t4, \$t6

TARGET:
sub \$t5, \$t0, \$t5

inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
beq	F	D	X	M	W								
sub		F	D	X	M	W							

Predict taken (correct!)



Branch prediction

- Dynamic branch direction prediction:
 - *Predict* branch outcome (T or N) in IF stage
 - Assume branches resolved in ID stage

beq \$t0, \$t1, TARGET
add \$t5, \$t4, \$t6

TARGET:
sub \$t5, \$t0, \$t5

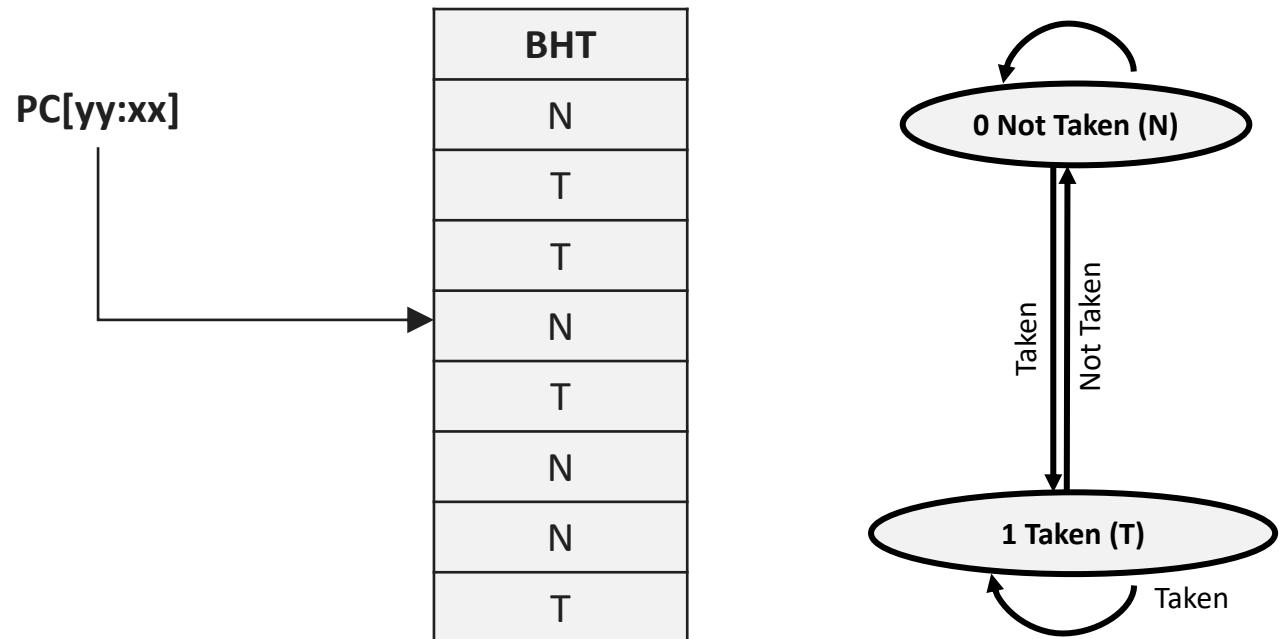
inst\cycle	1	2	3	4	5	6	7	8	9	10	11	12	13
beq	F	D	X	M	W								
add			F	=									
sub				F	D	X	M	W					

Predict not taken (incorrect)

Branch history table

□ 1-bit prediction

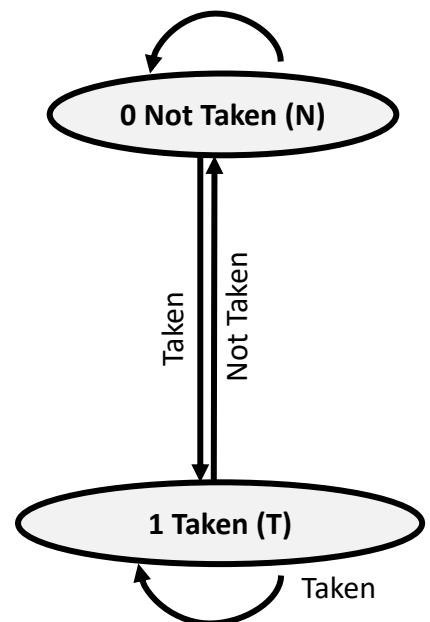
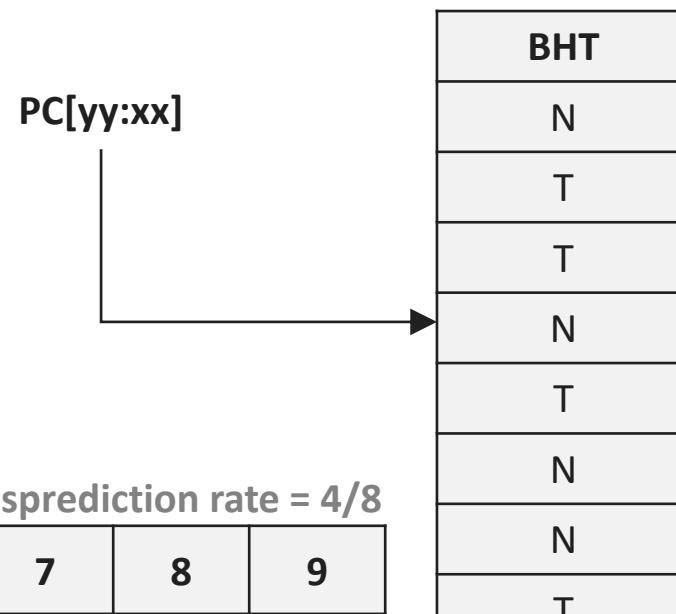
- Why use PC for addressing?
- Why drop PC's least significant bit(s)?



Branch history table

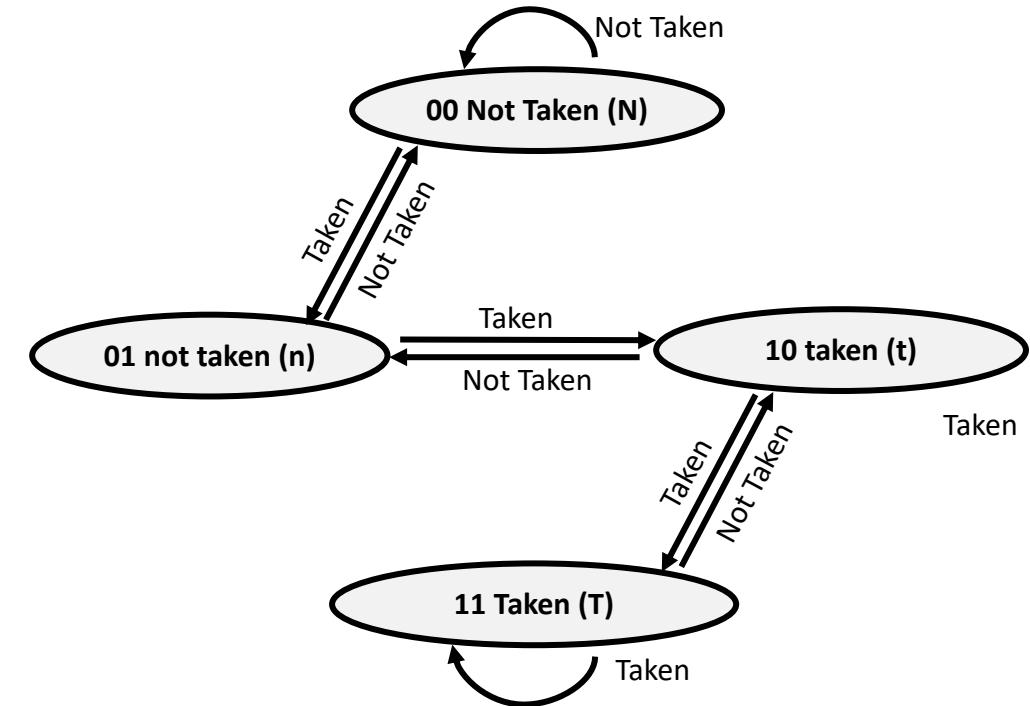
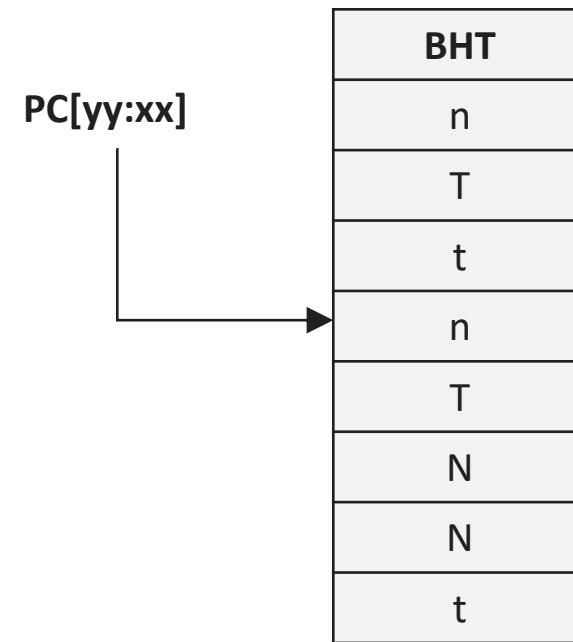
□ 1-bit prediction

- Why use PC for addressing?
- Why drop PC's least significant bit(s)?



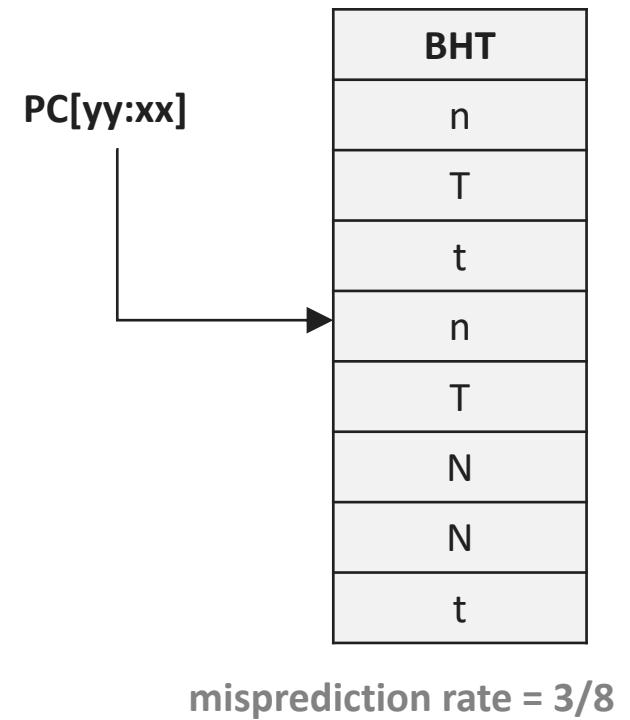
Branch history table

- ☐ 2-bit prediction

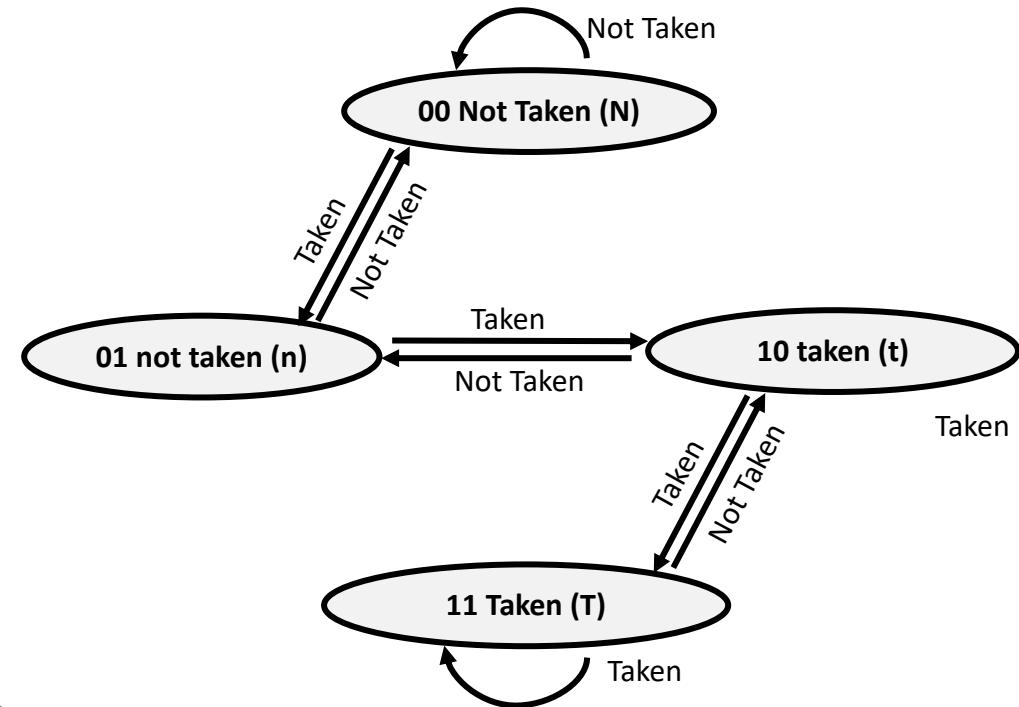


Branch history table

2-bit prediction



LOOP iteration	1	2	3	4	6	7	8	9
prediction	n	t	T	T	t	T	T	T
outcome	T	T	T	N	T	T	T	N





Exercise

❑ Consider five-stage pipeline

- Assume accounting for data hazards, average CPI = 1.2
- Assume branches resolved in ID stage
- Assume 30% of instructions are branches: 60% are taken

Q. Calculate average CPI and throughput

a. Predict-not-taken with 300us clock period?

b. 90%-accurate predictor with 325us clock period?



Midterm 1 outline



Questions?
